

Rahul Naidu Siriporam

+91 95899 32860

rahulnaidusiriporam@gmail.com

linkedin.com/in/rahul-naidu-siriporam-15472116b

x.com/rahulwlw

github.com/rnaidu-parallel

blog.neuromancer.in

SUMMARY

Founding software engineer with 5 years of end-to-end product experience, architecting and shipping backend platforms, data pipelines, and LLM systems from scratch into production. At Canvas AI, has owned consumer-health AI products, agentic chat platforms, vision pipelines, Text-to-SQL bots, WhatsApp-native AI companions, and a ClickHouse-backed analytics warehouse. Comfortable leading small teams, owning problems end-to-end with minimal direction, combining strong engineering execution with product intuition, and accelerating delivery using agentic coding tools (Claude Code, Codex). Publishes technical articles at blog.neuromancer.in on agentic AI systems, including local LLM deployments and Hermes-based agents on consumer hardware.

SKILLS

Languages: Python, TypeScript, SQL

AI/ML: Multi-Agent Systems, Agent Evals, RAG, GraphRAG, Fine-tuning, Prompt Engineering, Prompt Caching, Transformers, PyTorch

Backend: FastAPI, Hono, Flask, Next.js, Kafka, Debezium, Airflow, Drizzle

Infra: Docker, Docker Swarm, Kubernetes, GKE, KEDA, Pulumi, Datadog, CI/CD

Data: PostgreSQL, MongoDB, ClickHouse, Redis, SQLite

Cloud: AWS, GCP, Azure, ECS, Vertex AI, Azure AI

EXPERIENCE

Canvas AI (previously Apperture / Parallel) — Founding Software Engineer Jul 2022 – Present | Remote

- Built Jarvis, a WhatsApp-native agentic investing companion with a two-tier memory architecture (a bounded, structured memory card plus recallable message search) kept current by mid-session context compaction, session-close summary generation, and a two-phase memory-update agent, plus a scheduled proactive-outreach scanner that follows up on each user's open loops. Both products run on ECS with per-user queueing, tenant-isolated Postgres, and prompt caching for sub-linear token spend on long-running conversations, with an LLM judge and eval harness scoring every turn (structured prompt and response logging, latency and accuracy dashboards).
- Architected a consumer-health AI assistant for diabetes care, leading a team of 3 engineers, streaming to thousands of users daily: evolved it from RAG to GraphRAG to a QLoRA-fine-tuned model over the SugarFit corpus, migrated from Azure OpenAI and Prompt Flow to Vertex AI, and re-architected from a single agent into a multi-agent system on Agno (Gemini 2.5 Pro/Flash), supporting meal planning, symptom triage, CGM analysis, lab-report parsing (pdfplumber + LLM extraction), and slot-level "Day at a Glance" insights.
- Shipped a clinician copilot (auto-suggested diagnoses and prescriptions), a coach pre-call summarizer over Gemini audio understanding, and an automated webinar Q&A responder, on GKE with KEDA autoscaling and Datadog, tuned for latency and cost via prompt engineering and per-task model routing.
- Replicated the CGMformer paper to deploy a hybrid Transformer-LSTM glucose forecasting model on Vertex AI, achieving 94% of 2-hour-ahead predictions within Clarke Error Grid Zone A; and a vision meal-logging pipeline on Gemini 2.5 Flash multimodal processing 10K+ images per day at p90 ~6s end-to-end, lifting in-app meal-logging adoption from 24% to 38% (+58%).
- Built IRA, a Next.js agentic investment-research assistant (in private beta) on the Vercel AI SDK, whose skill-loading ToolLoopAgent drives a deterministic multi-factor scoring engine and an interactive analysis dashboard, with human-in-the-loop clarification.
- Built the core of Apperture, a B2B product-analytics platform serving real customers: funnel, retention, and cohort analysis, user segmentation, and custom metrics on a dynamic query builder over ClickHouse, plus 8+ data connectors (GA, Mixpanel, Amplitude, PostHog, and major SQL and NoSQL databases) unifying ingestion of tens of millions of events per day, on a production stack of Debezium CDC, Kafka streaming, Airflow orchestration, and AWS microservices on Docker Swarm with full CI/CD and monitoring.

Quantile Analytics — Quantitative Analyst Jul 2021 – Jul 2022 | Bengaluru

- Built automated pipelines to track 20+ global indices and generate rebalancing signals, replacing manual workflows and enabling systematic strategies on US and international equities.
- Built an NLP parser for SEC filings (float and share-data extraction), exchange-data scrapers that fed index-rebalancing models, and Flask and Angular internal tools that improved data accessibility and streamlined the research team's workflows.

EDUCATION

Indian Institute of Technology, Indore — B.Tech 2017 – 2021

ACHIEVEMENTS

- JEE Main 2017: AIR 3,289 (top 0.27% among 1.2M candidates). JEE Advanced 2017: AIR 4,473. CAT 2021: 99.19 percentile.
- Bronze Medal, BitGrit Data Science Contest, Inter-IIT Tech Meet, IIT Roorkee 2019.